

✓ 3MTT Capstone Project




Introduction

Riverine communities need flood-risk awareness in order to prepare for and avoid damages caused by flooding. This model classifies a community's flood risk from environmental features. Data were captured at 722 locations within the project area. Thus, community members can determine how exposed they are to flood risk by simply imputing their location to the model.

The model is built around Ero river catchment area. Ero River is in Moba and Ilejemeje LGAs in Ekiti State serving at least seventeen communities with water for irrigation and domestic use. A dam is built on the Ero river at Ikun-Ekiti. But the river causes flooding during raining seasons, and as a result of the operations of the dam.

Ero River Catchment Area



Ero River is in Moba and Ilejemeje LGAs serving at least seventeen communities in Ekiti State

✓ Preparing the Dataset

✓ Data Source

In 2025, two PhD candidates, Andrew ToluTaiwo and Thomas Olonisaye, of the University of Lagos did a study on flood susceptibility mapping in Ero River catchment area in Ekiti State, Nigeria. They collected the Flood Predictors dataset used in this project.

✓ Flood Predictors

Flood Predictors are environmental features, topographic and hydrologic variables, that influence the occurrence or magnitude of flooding incidents.

- **lon: Longitude.** Longitude coordinate of a location
- **lat: Latitude.** Latitude coordinate of a location
- **dist_ri: Distice from River.** Distance in metres from the nearest river
- **elev: Elevation.** Height above mean sea level of the location
- **NDVI: Normalized Difference Vegetation Index.** A measure of occurrence of vegetation in the location.
- **NDWI: Normalized Difference Water Index.** A measure of the wetness of the soil.
- **rain: Rainfall.** Amount of rainfall in the region

- **NDBI: Normalized Difference Building Index.** Differentiates between built-up areas and non built-up areas.
- **LULC: Land Use Land Cover.** Land Use refers to how humans use land for various purposes. Land Cover refers to the materials that cover a part of the earth's surface. 0-Water body, 1-Farmland, 2-Forest, 3-Builtup, 4-Bareground
- **rock_por: rock porosity.** Form and porousness of sub-surface rock.
- **rock_perm: Rock Permeability.** Rock permeability affects runoff.
- **Soil_por: Soil Porosity.** Form and porousness of soil.
- **soil_perm: Soil Permeability.** Soil permeability affects runoff.

flood_risk: Flood Risk Class. A location's susceptibility to flooding. This is the target variable.

```
import pandas as pd
import numpy as np
import random
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# supress warnings
import warnings
warnings.filterwarnings('ignore')
```

```
## Load the data
flood_data = pd.read_csv('flood_prediction_data2.csv')
flood_data = flood_data.drop(columns=['Unnamed: 0'])
```

✓ A little bit of data exploration and cleaning

```
## Explore the first 5 rows
flood_data.head(5)
```

	lon	lat	dist_ri	elev	NDVI	NDWI	rain	NDBI	LULC	rock_por	rock_perm	soil_por	soil_perm
0	5.1572	8.0133	5	517	0.55	-0.66	259.75	-0.07	1	0.05	1.000000e-16	0.55	0.001
1	5.1707	8.0147	5	519	0.48	-0.55	255.58	0.01	1	0.05	1.000000e-16	0.55	0.001
2	5.1649	8.0072	5	531	0.47	-0.68	254.91	-0.01	1	0.05	1.000000e-16	0.55	0.001
3	5.1810	8.0033	5	526	0.48	-0.60	254.00	0.00	1	0.05	1.000000e-16	0.55	0.001
4	5.1832	8.0108	5	514	0.42	-0.63	254.00	0.04	4	0.05	1.000000e-16	0.55	0.001

Next steps: [Generate code with flood_data](#) [New interactive sheet](#)

✓ Missing values

```
# Where are there missing values?
missing_vals = flood_data.isna().sum()
```

```
missing_vals.index[missing_vals>0]
missing_vals
```

```

      0
lon    0
lat    0
dist_ri 0
elev    0
NDVI    0
NDWI    0
rain    0
NDBI    0
LULC    0
rock_por 0
rock_perm 0
soil_por 0
soil_perm 0
flood_risk 0

```

dtype: int64

```
# print important statistics
flood_data.describe()
```

	lon	lat	dist_ri	elev	NDVI	NDWI	rain	NDBI	
count	722.000000	722.000000	722.000000	722.000000	722.000000	722.000000	722.000000	722.000000	722.000000
mean	5.152463	7.956060	3.130194	546.055402	0.569889	-0.618352	306.156884	-0.077271	3.000000
std	0.066029	0.033578	1.634307	21.812013	0.135047	0.064207	33.879976	0.065769	1.290000
min	5.030230	7.872770	1.000000	477.000000	0.130000	-0.740000	0.000000	-0.200000	0.000000
25%	5.090160	7.933805	1.000000	529.000000	0.490000	-0.670000	283.990000	-0.130000	2.250000
50%	5.153855	7.956460	4.000000	548.000000	0.565000	-0.620000	312.355000	-0.080000	4.000000
75%	5.209290	7.981400	5.000000	562.000000	0.670000	-0.570000	332.767500	-0.030000	4.000000
max	5.269120	8.029350	5.000000	603.000000	0.820000	-0.380000	361.240000	0.090000	4.000000

Investigating Feature Correlation with `suscep` for Data Leakage

```


correlation_matrix = flood_data.corr()
suscep_correlations = correlation_matrix['flood_risk'].sort_values(ascending=False)
display(suscep_correlations)

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix[['flood_risk']].sort_values(by='flood_risk', ascending=False), annot=True)
plt.title('Feature Correlation with Flood Risk (flood_risk)')
plt.show()

```

	flood_risk
flood_risk	1.000000
dist_ri	0.780005
soil_por	0.492618
NDBI	0.448793
lon	0.423115
soil_perm	0.161554
lat	0.144471
rock_por	-0.012874
rock_perm	-0.019067
NDWI	-0.060620
LULC	-0.277017
rain	-0.283517
elev	-0.400199
NDVI	-0.543188

dtype: float64



✓ Training, Validation and Testing Splits

We often split the dataset into three different chunks:

1. **Training data:** the data we build our models on.
2. **Validation data:** the data we tune hyperparameters on.
3. **Testing data:** the data we use to obtain a final estimate of performance.

The validation set is like an 'intermediate' estimate of performance. If we didn't use the validation set for this purpose, the only way of selecting the best model from a model class would be to look at its performance on the training set or the testing set.

`train_test_split` in scikit-learn provides an easy way of creating these data partitions. Let's start by just creating training/testing sets. We'll get to parameter tuning in a minute.

```
from sklearn.model_selection import train_test_split

X = flood_data.drop(columns=['flood_risk'], axis=1)
y = flood_data['flood_risk']

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    train_size = 0.70, random_state = 1)

print("Training Size = %d" % X_train.shape[0]) # To print integers, we can pass '%d' into our print sta
print("Testing Size = %d" % X_test.shape[0])

Training Size = 505
Testing Size = 217
```

The `sklearn.preprocessing` facilitates the transition from 'raw' data to a standard format for algorithm training.

Among such algorithms, we find:

Standardization

Standardization is the first type of transformation to consider, as a large number of machine learning/statistical algorithms assume that the data they operate on is normally distributed. Standardization is expressed as:

$$\frac{x_i - \text{mean}(x)}{\text{std}(x)}$$

In practice, the shape of the distribution to work with is ignored and the transformation is simply performed by removing the mean and scaling by the standard deviation.

Note: This is not recommended if the histogram of the target variable is far from Gaussian. In such cases, alternative transformations will be reviewed.

Min-Max Scaling

A good alternative to the previous method is range scaling, which is expressed as:

$$\frac{x_i - \min(x)}{\max(x) - \min(x)}$$

for x being the column to treat, and x_i the element to transform. This transformation allows the data to range between 0 and 1 and can be used when the distribution of the data does not meet the normality assumption. Note that this transformer is affected by the presence of outliers.

```
### The "manual" way ###
```

```
# Normalization
X_train_normalized = (X_train - X_train.min())/(X_train.max()-X_train.min())
X_test_normalized = (X_test - X_train.min())/(X_train.max()-X_train.min())

# Standardization
X_train_standardized = (X_train - X_train.mean())/X_train.std()
X_test_standardized = (X_test - X_train.mean())/X_train.std()
```

✓ Logistic Regression

Logistic regression is part of the class of generalized linear models (GLMs), which build directly on top of linear regression. These models take the linear fit and map it through a non-linear function.

Specifically, logistic regression is used when the dependent variable is **binary** (e.g. true/false, yes/no, 0/1) and the goal is to find the relationship between one or more independent variables and the probability of the dependent variable being true or 1.

For logistic regression, the function $f(\cdot)$ is given by $f(z) = \frac{1}{1 + \exp(-z)}$ (the *logistic function* - also called the *sigmoid function*).

Where:

- $f(z)$ is the probability of the dependent variable being true or 1
- z is the linear combination of the independent variables and the coefficients of the model:

$$z = b_0 + b_1 x_1 + b_2 x_2 + \dots + b_n x_n$$

We use a grid search to tune regularization parameters:

1. Define a parameter grid
2. Initialize a GridSearchCV object with the tuning grid and other fixed parameters
3. Fit the model!

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV

# Define the grid that we want to search over
param_grid = {'C': np.arange(0.001, 1, 0.05),
              'penalty': ['l2', 'l1'],
              'solver': ['liblinear']}

# Define the parameters for the model
gs = GridSearchCV(LogisticRegression(random_state=42, max_iter = 1000),
                  return_train_score=True,
                  param_grid=param_grid,
                  scoring='roc_auc',
                  cv=5, verbose = 0)

## Fit the model
random.seed(42)
gs.fit(X_train_standardized, y_train)
m_lr = gs.best_estimator_
print("Best parameters: ", gs.best_params_)

Best parameters: {'C': np.float64(0.15100000000000002), 'penalty': 'l1', 'solver': 'liblinear'}
```

```
# Let's look at the coefficients
coef_output = pd.DataFrame({'feature': X_train_standardized.columns,
                           'coefficient': m_lr.coef_[0]})
```

```
print("Number of zeros: %d" % (coef_output.query('coefficient == 0').shape[0]))
coef_output.sort_values('coefficient')
```

Number of zeros: 8

	feature	coefficient	
4	NDVI	-1.382105	
3	elev	-0.999582	
5	NDWI	-0.504173	
6	rain	-0.169481	
0	lon	0.000000	
1	lat	0.000000	
7	NDBI	0.000000	
8	LULC	0.000000	
12	soil_perm	0.000000	
9	rock_por	0.000000	
10	rock_perm	0.000000	
11	soil_por	0.000000	
2	dist_ri	2.152703	

Model Evaluation

We can now explore out-of-sample performance. Ultimately, we want to predict whether or not a location is flooded. Logistic regression gives a probability that each location is flooded. This gives us flexibility in the way we predict. The most natural thing would be to predict that any location with predicted probability above 0.5 is flooded, and any location with predicted probability below 0.5 is not flooded. But if we want to be certain, may want to use a cutoff value of 0.9 rather than 0.5. In fact, we could choose any cutoff value and have a corresponding prediction model.

This is where AUC comes in. For every cutoff, we'll plot the *false positive rate* (FPR) against the *true positive rate* (TPR) and then take the area under this curve.

Also known as **Sensitivity** or **Recall**, TPR measures the proportion of actual positives that are correctly identified by the model.

$$TPR = \frac{TP}{TP+FN}$$

Where:

- *TP* (True Positives) = Correctly predicted positive cases.
- *FN* (False Negatives) = Actual positive cases that were incorrectly predicted as negative.

FPR measures the proportion of actual negatives that are incorrectly identified as positive by the model.

$$FPR = \frac{FP}{FP+TN}$$

Where:

- *FP* (False Positives) = Incorrectly predicted positive cases.
- *TN* (True Negatives) = Correctly predicted negative cases.

High TPR means the model correctly identifies most of the positive cases. **Low FPR** means the model makes fewer false alarms by incorrectly classifying negatives as positives.


```
train_pred = m_lr.predict_proba(X_train_standardized)[: ,1]
test_pred = m_lr.predict_proba(X_test_standardized)[: ,1]

print("Train AUC: %.3f" % metrics.roc_auc_score(y_train, train_pred))
print("Test AUC: %.3f" % metrics.roc_auc_score(y_test, test_pred))

train_fpr, train_tpr, _ = metrics.roc_curve(y_train, train_pred)
test_fpr, test_tpr, _ = metrics.roc_curve(y_test, test_pred)

# plot the roc curve for the model
plt.plot(train_fpr, train_tpr, linestyle='--', label='Train')
plt.plot(test_fpr, test_tpr, marker='.', label='Test')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
# show the legend
plt.legend()
# show the plot
plt.show()
```

Train AUC: 0.980
Test AUC: 0.987



✓ Confusion Matrix Analysis


```
t = 0.5
print("Train Accuracy: %.3f" % metrics.accuracy_score(y_train, train_pred > t))
print("Test Accuracy: %.3f" % metrics.accuracy_score(y_test, test_pred > t))

# If we wanted to get all of the elements of the confusion matrix, we could pull them like this:
tn, fp, fn, tp = metrics.confusion_matrix(y_test, test_pred > t).ravel()

# Plot the confusion matrix on the test set
cm = metrics.confusion_matrix(y_test, test_pred > t)
```

```
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

Train Accuracy: 0.931
Test Accuracy: 0.940



From these numbers, we can see that the model is performing quite well, with a higher number of true positives and true negatives. The number of false negatives (4) is lower than false positives (9), which means the model is less likely to miss actual flood-susceptible locations, which is often a critical aspect in flood prediction tasks.

```
np.mean(1-y_train)
```

```
np.float64(0.5544554455445545)
```

Note: The code `np.mean(1-y_train)` calculates the proportion of instances in the training set (`y_train`) where the flood susceptibility (`suscep`) is 0. In other words, it tells us the percentage of training samples that are not susceptible to flooding.

The result, 0.5544554455445545 (approximately 55%), indicates that about 55% of the training data points represent locations that are not susceptible to flooding. This value can also be considered the baseline accuracy if a model were to simply predict 'no flood' for every instance, without any learning.

✓ K-Fold Cross-Validation for Robust Accuracy Estimate

```
from sklearn.model_selection import cross_val_score
```

```
# Perform 5-fold cross-validation on the logistic regression model
cv_scores = cross_val_score(m_lr, X, y, cv=5, scoring='accuracy')

print(f"Cross-validation accuracy scores: {cv_scores}")
print(f"Mean cross-validation accuracy: {np.mean(cv_scores):.3f}")
print(f"Standard deviation of cross-validation accuracy: {np.std(cv_scores):.3f}")
```

Cross-validation accuracy scores: [1. 0.90344828 0.97916667 0.79861111 0.70138889]
Mean cross-validation accuracy: 0.877
Standard deviation of cross-validation accuracy: 0.112

Mean cross-validation accuracy: 0.877

This is the average of the accuracy scores from all 5 folds. It gives an overall, more generalized estimate of how well the model is expected to perform on unseen data. An average of 87.7% is quite good.

Standard deviation of cross-validation accuracy: 0.112. This tells how much the individual fold scores deviate from the mean. A higher standard deviation (like 0.112 in this case, relative to the mean) suggests that the model's performance is somewhat inconsistent across different data subsets.

✓ Flood Risk Categories

```
import numpy as np

# Standardize the entire X dataset using the mean and std from X_train
X_standardized = (X - X_train.mean()) / X_train.std()

# Ensure X is the full feature set as used for training
# Get probability predictions from the best logistic regression model for the full dataset
flood_probabilities = m_lr.predict_proba(X_standardized)[: , 1]

# Define the risk classification function
def classify_flood_risk(probability):
    if 0.00 <= probability <= 0.20:
        return 'Low Flood-Risk'
    elif 0.21 <= probability <= 0.45:
        return 'Moderate Flood-Risk'
    elif 0.46 <= probability <= 0.70:
        return 'High Flood-Risk'
    elif 0.71 <= probability <= 1.00:
        return 'Very High Flood-Risk'
    else:
        return 'Unknown'

# Create a copy of the original flood_data to add new columns
flood_data_classified = flood_data.copy()

# Add probability and risk classification to the DataFrame
flood_data_classified['flood_probability'] = flood_probabilities
flood_data_classified['flood_risk_category'] = flood_data_classified['flood_probability'].apply(classify_flood_risk)

# Display the first few rows with the new columns
print("Flood data with risk classification:")
display(flood_data_classified)

# Save the results to a CSV file
output_csv_path = 'flood_risk_classified_data.csv'
flood_data_classified.to_csv(output_csv_path, index=False)
print(f"Classified flood risk data saved to '{output_csv_path}'")
```

Flood data with risk classification:

	lon	lat	dist_ri	elev	NDVI	NDWI	rain	NDBI	LULC	rock_por	rock_perm	soil_por	soil_
0	5.15720	8.01330	5	517	0.55	-0.66	259.75	-0.07	1	0.0500	1.000000e-16	0.55	0.00
1	5.17070	8.01470	5	519	0.48	-0.55	255.58	0.01	1	0.0500	1.000000e-16	0.55	0.00
2	5.16490	8.00720	5	531	0.47	-0.68	254.91	-0.01	1	0.0500	1.000000e-16	0.55	0.00
3	5.18100	8.00330	5	526	0.48	-0.60	254.00	0.00	1	0.0500	1.000000e-16	0.55	0.00
4	5.18320	8.01080	5	514	0.42	-0.63	254.00	0.04	4	0.0500	1.000000e-16	0.55	0.00
...	...	...	...	...	...	...	...	...	...	...	...	...	...
717	5.13083	7.94538	4	531	0.66	-0.60	333.05	-0.14	4	0.0046	3.000000e-03	0.58	0.00
718	5.12168	7.94011	3	533	0.58	-0.56	333.22	-0.11	4	0.0046	3.000000e-03	0.58	0.00
719	5.13720	7.93401	4	549	0.61	-0.60	333.00	-0.10	4	0.0046	3.000000e-03	0.58	0.00
720	5.21341	7.90935	4	525	0.64	-0.61	361.00	-0.08	3	0.0500	1.000000e-16	0.58	0.00
721	5.24833	7.90519	3	525	0.50	-0.54	354.10	-0.10	4	0.0500	1.000000e-16	0.58	0.00

722 rows × 16 columns

Classified flood risk data saved to 'flood_risk_classified_data.csv'

Next steps: [Generate code with flood_data_classified](#) [New interactive sheet](#)

✓ Spatial Distribution of flood-risk categories.

```

import matplotlib.pyplot as plt
import seaborn as sns

# Define a custom color palette to ensure distinct colors for each risk category
# and order them logically (e.g., green for low, red for very high)
color_palette = {
    'Low Flood-Risk': 'green',
    'Moderate Flood-Risk': 'blue',
    'High Flood-Risk': 'orange',
    'Very High Flood-Risk': 'red'
}

# Ensure the categories are ordered for consistent plotting and legend
risk_order = ['Low Flood-Risk', 'Moderate Flood-Risk', 'High Flood-Risk', 'Very High Flood-Risk']
flood_data_classified['flood_risk_category'] = pd.Categorical(
    flood_data_classified['flood_risk_category'], categories=risk_order, ordered=True
)

plt.figure(figsize=(12, 10))
sns.scatterplot(
    data=flood_data_classified,
    x='lon',
    y='lat',


```

```

hue='flood_risk_category',
palette=color_palette,
s=50, # Size of the points
alpha=0.7, # Transparency of the points
edgecolor='w', # White edge for better distinction
linewidth=0.5 # Line width of the edge
)

plt.title('Spatial Distribution of Flood Risk Categories')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.legend(title='Flood Risk Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

```



✓ Number of locations in the floo_risk categories

```

print('Number of locations in each flood risk category:')
display(flood_data_classified['flood_risk_category'].value_counts().sort_index())

```

Number of locations in each flood risk category:

	count
flood_risk_category	
Low Flood-Risk	314
Moderate Flood-Risk	76
High Flood-Risk	44
Very High Flood-Risk	277

dtype: int64

Double-click (or enter) to edit

Distribution Analysis

This distribution of flood risk categories provides valuable insights into the flood risk of the evaluated locations:

- **Low Flood-Risk:** These locations have a low probability of flooding, suggesting they are relatively safe. For these areas, standard preventative measures might suffice, and resources can be allocated elsewhere.
- **Moderate Flood-Risk:** These areas show a noticeable but not extreme probability of flooding. They might require regular monitoring and basic flood preparedness plans. Future development in these areas should consider drainage and infrastructure to mitigate potential risks.
- **High Flood-Risk:** Locations in this category have a significant chance of experiencing floods. These areas warrant more immediate attention, including early warning systems, robust flood defenses, and strict building codes. Community awareness and evacuation plans are crucial here.
- **Very High Flood-Risk:** These are the most vulnerable locations, with a very high likelihood of flooding. Urgent and comprehensive flood mitigation strategies are needed, potentially including relocation considerations, major infrastructure projects, and stringent land-use planning. These areas could be prioritized for immediate interventions and resource allocation.

This classification is a powerful tool for urban planners, emergency services, and residents to make informed decisions regarding development, resource allocation, and safety measures in the Ero River catchment area.

Conclusion

This project successfully developed and evaluated a Logistic Regression model for classifying flood risk in the Ero River catchment area based on various environmental features. By leveraging detailed geospatial and environmental data, the model provides valuable insights into the flood susceptibility of 722 locations.

Implications:

The generated flood risk map and classifications can serve as a crucial tool for:

- **Urban Planning:** Informing decisions on new developments and infrastructure projects to minimize exposure to flood hazards.
- **Emergency Management:** Guiding the deployment of resources, establishment of early warning systems, and planning of evacuation routes.
- **Community Awareness:** Educating residents about their specific flood risk and promoting preparedness measures.